

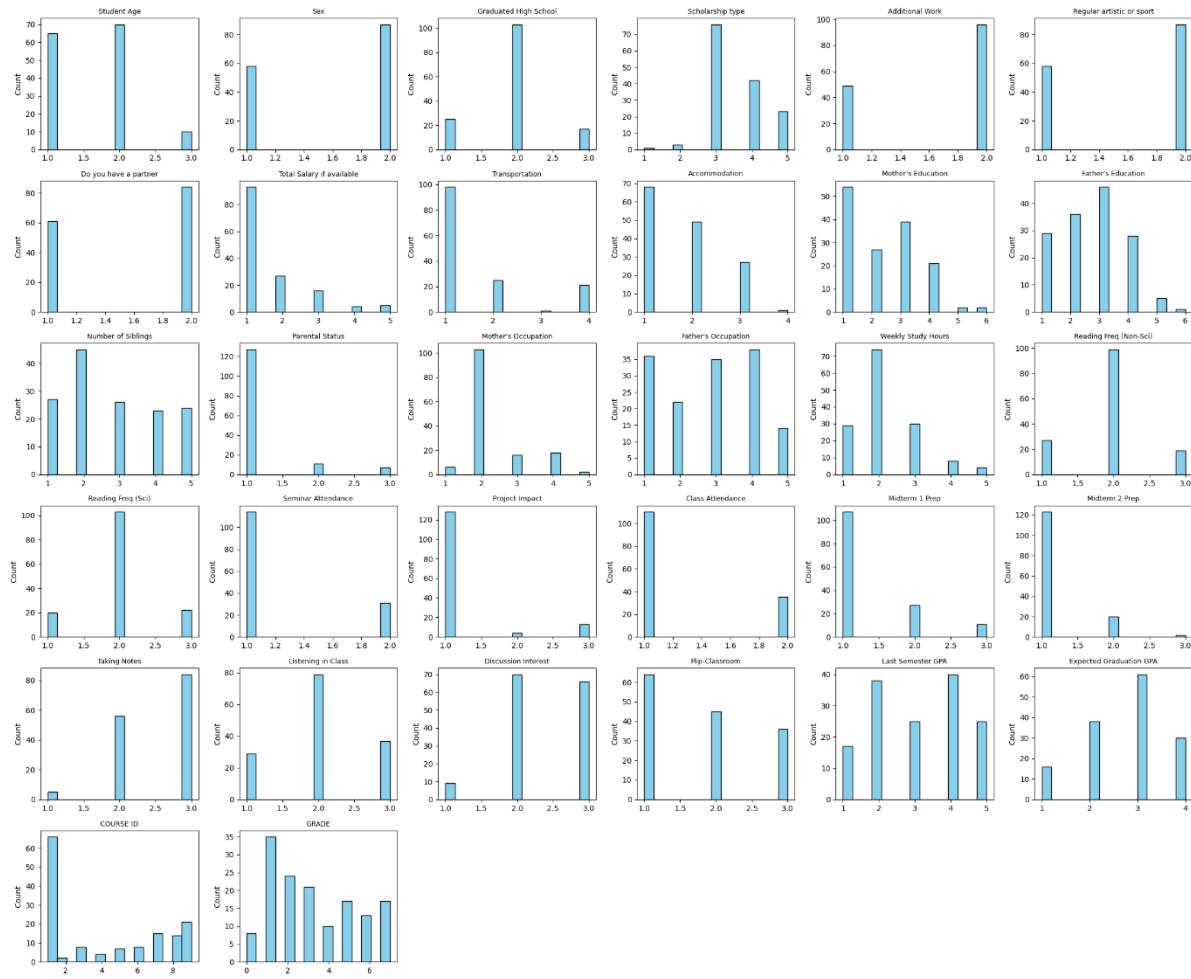
Student Performance Analysis: Guide to Diagrams

Dataset Overview

Before diving into the visualizations, here is a simple breakdown of the data we analysed.

- **What is it?** A collection of performance records from a higher education evaluation system.
- **Size:** We have data for **145 students** (rows) covering **33 different attributes** (columns).
- **Key Data Points:**
 - **Student ID:** A unique label for each student (no duplicates).
 - **Questions 1–30:** These are survey questions about habits, family, and lifestyle. Students answered on a scale of **1 to 5**.
 - **Course ID:** Indicates which of the 9 courses the student took (numbered 1–9).
 - **Final Grade:** This is the score we are trying to predict. It ranges from **1 (Lowest)** to **7 (Highest)**.
- **Quality:** The dataset is very clean, meaning there were no missing values or blank spots we had to fix.
- **Why the Dataset Is Useful**
 - The dataset is suitable for:
 - Predicting student grades
 - Identifying factors that influence academic outcomes
 - Building ML models such as Logistic Regression or Random Forest
 - Understanding student performance across courses

1. Histograms (Distribution of Data)



What is this?

Think of this as a "headcount" for every question in the survey.

- **The X-axis (bottom)** shows the possible answers (e.g., 1 to 5).
- **The Y-axis (side)** shows how many students picked that answer.

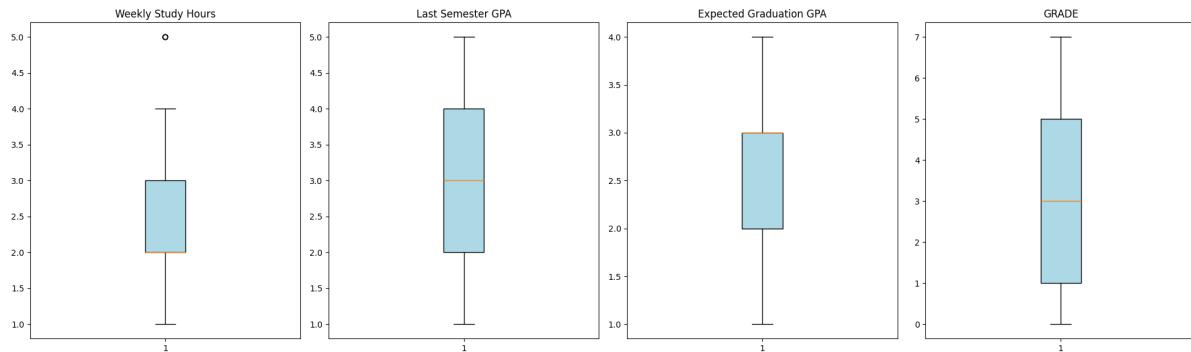
How we made it:

We took every column in the spreadsheet and asked Python to count the frequency of each number.

What it tells us:

- Most of our data is "categorical," meaning students chose from specific options (like "Yes/No" or "1-5").
- You can see gaps between the bars, which confirms the data isn't continuous (like exact height or weight) but discrete (like multiple-choice answers).

2. Boxplots (Outlier Detection)



What is this?

A boxplot is a way to look at the "spread" of grades or study hours.

- **The Box:** Represents the "middle 50%" of students. Most people fall inside this box.
- **The Line in the Middle:** The median (the exact middle student).
- **The Whiskers (Lines sticking out):** The range of typical students.
- **The Circles (Dots):** These are "outliers"—students who are very different from the rest (e.g., someone studying 0 hours or getting a perfect Grade 7 when most get 3).

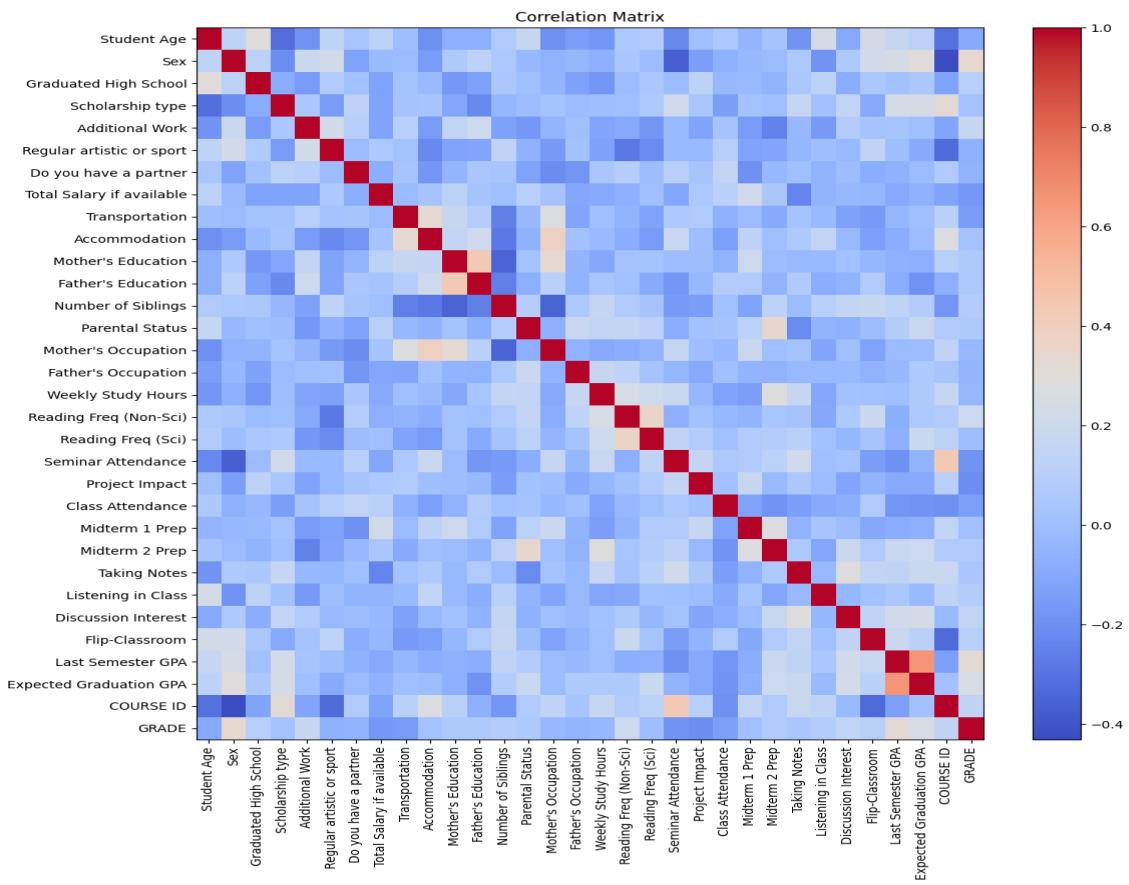
How we made it:

We selected key numeric columns (like Weekly Study Hours and GPA) and calculated their statistical range.

What it tells us:

- Our data is fairly stable.
- We can clearly see the range of grades (0 to 7), which is our target to predict.

3. Correlation Matrix (Heatmap)



What is this?

This is a grid that checks if two questions "move together."

- **Red / Warm Colours:** Positive relationship. (If X goes up, Y goes up).
- **Blue / Cool Colours:** Negative relationship. (If X goes up, Y goes down).
- **The Diagonal Line:** It is always dark red because every question matches itself perfectly.

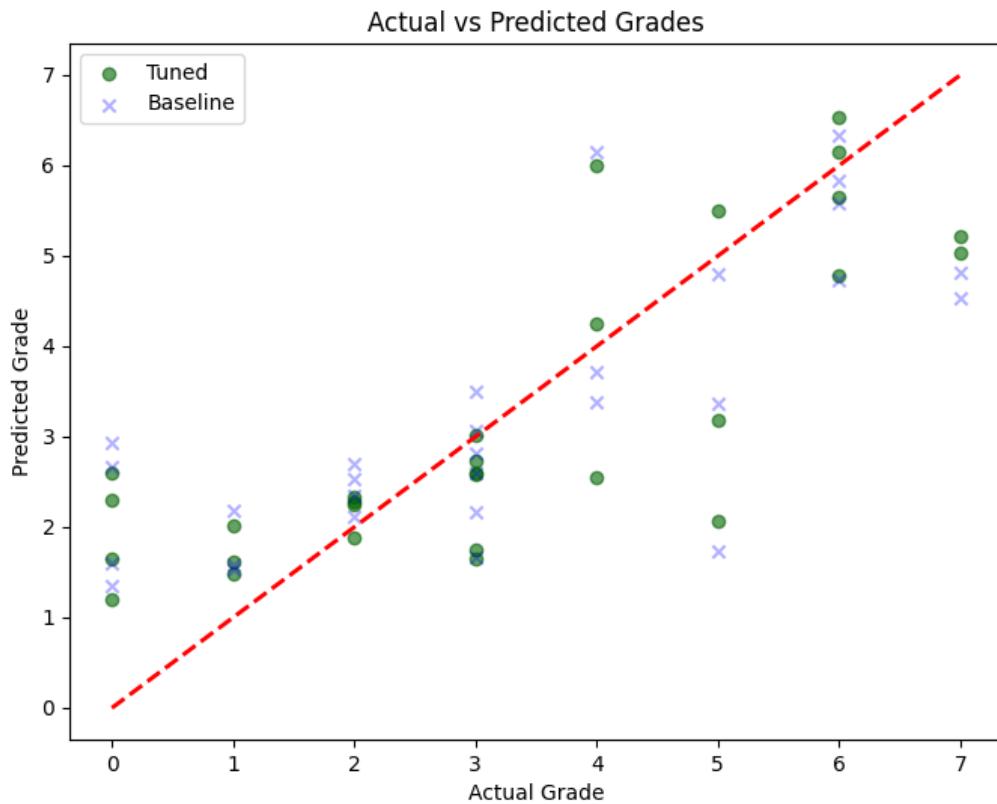
How we made it:

We used a mathematical formula called "Pearson Correlation" which gives a score between -1 and 1 for every pair of questions.

What it tells us:

- It helps us spot redundancy. If two questions are dark red, they probably ask the same thing.
- It helps us spot hints. We look for rows that are coloured (not white) relative to the GRADE column.

4. Actual vs. Predicted (Model Accuracy)



What is this?

This acts as a "report card" for our Artificial Intelligence model.

- **X-axis:** The grade the student *actually* got.
- **Y-axis:** The grade our model *guessed* they would get.
- **The Red Line:** This is the "Line of Perfection." If a dot lands on this line, the model guessed perfectly.

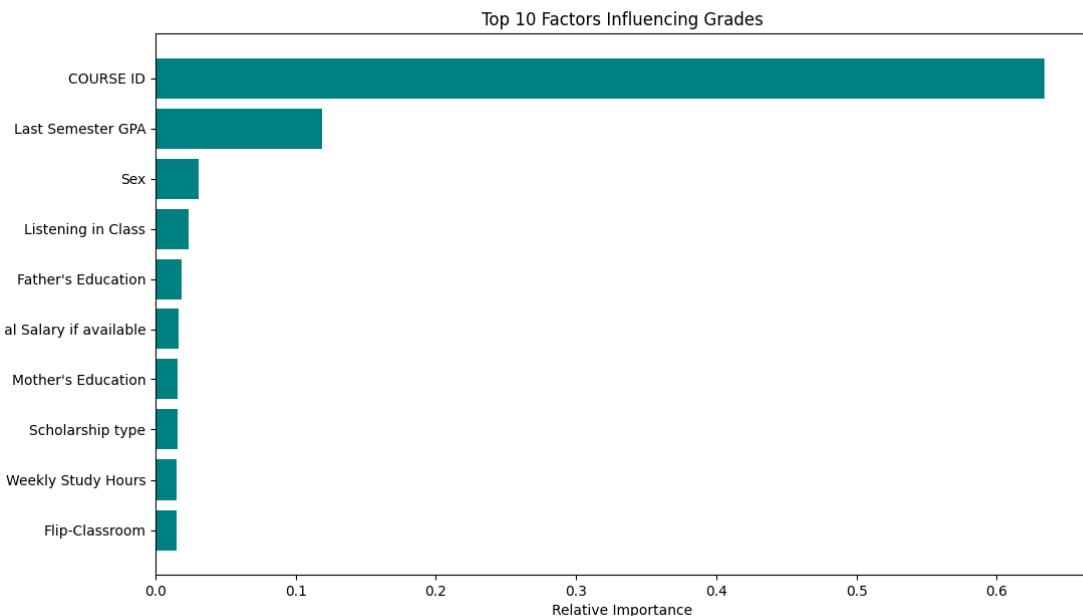
How we made it:

We hid 20% of the data (the "Test Set") from the model. After training, we asked the model to guess the grades for these hidden students and plotted the answers.

What it tells us:

- **Clustering:** The blue dots are grouped reasonably close to the red line. This means our model is generally accurate.
- **Error:** The distance of a dot from the red line represents the "error" (e.g., predicting a 4 when the student got a 6).

- **5. Feature Importance (Bar Chart)**



What is this?

This is a ranking of "What matters most?" It tells us which survey questions helped the AI predict the grade.

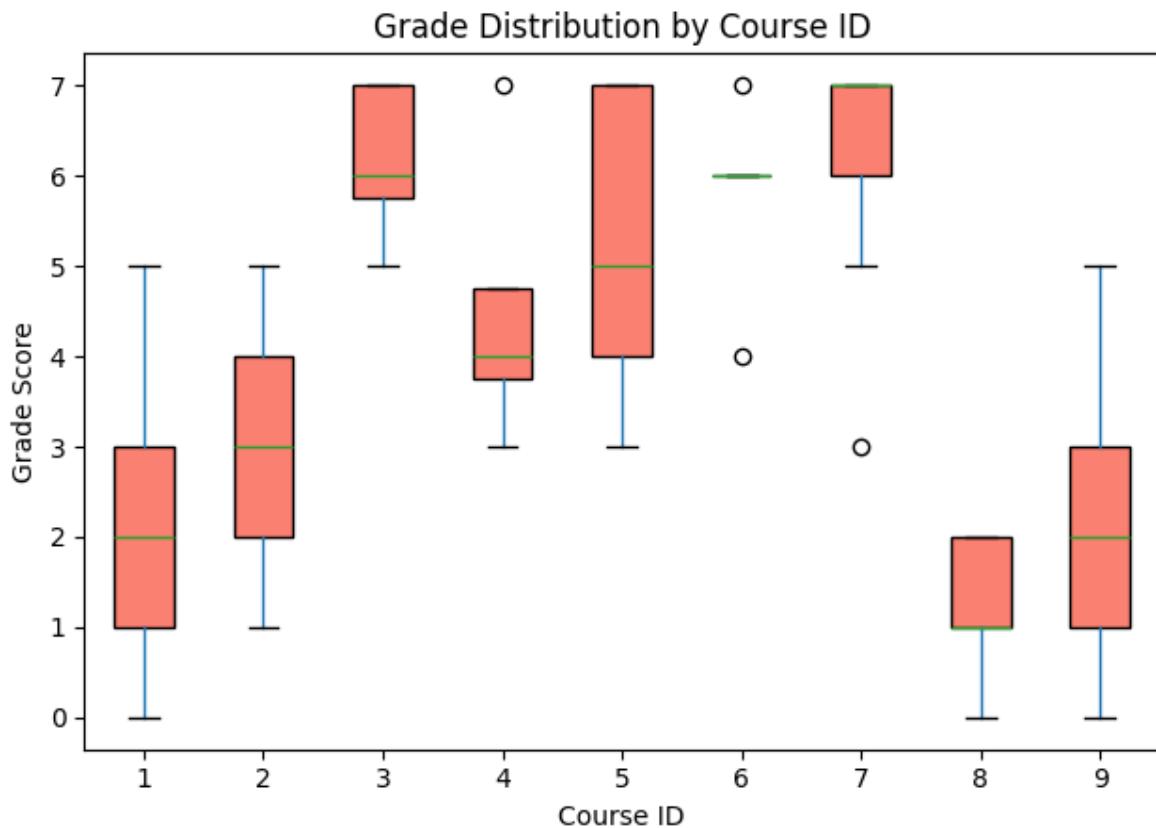
How we made it:

We asked our Random Forest model: "When you make decisions, which questions do you look at the most?"

What it tells us:

- **The #1 Driver:** COURSE ID is the biggest bar by far. This means the specific course a student takes matters more than their individual habits.
- **The #2 Driver:** Expected Graduation GPA. Students who set high goals perform better.

6. Course ID Impact (Grade by Course)



What is this?

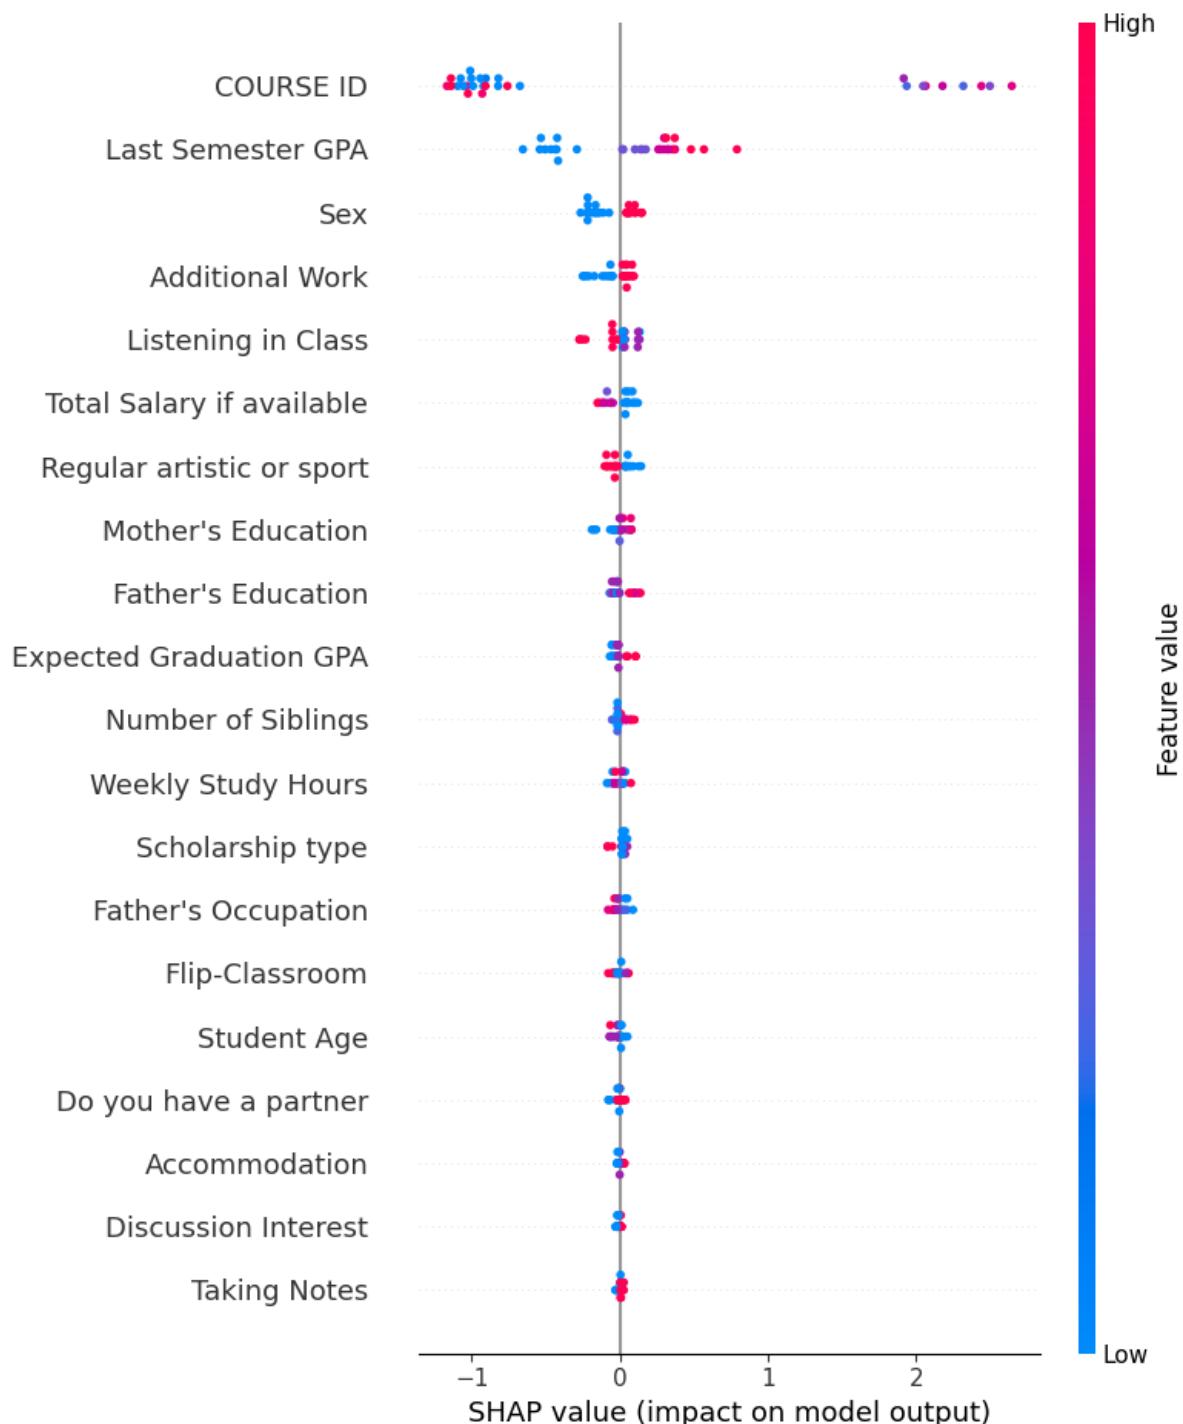
Since COURSE ID was the #1 predictor, we zoomed in to see why.

- Each box represents a different course (Course 1, Course 2, etc.).
- The position of the box shows the average grade in that course.

What it tells us:

- **Huge Variation:** Some courses (like Course 7) have very high grades (average ~6.0). Others (like Course 8) have very low grades (average ~1.4).
- **Conclusion:** The model relies on Course ID because some courses are simply "easier" or "harder" (or graded differently) than others.

7. SHAP Summary Plot (The "Why")



What is this?

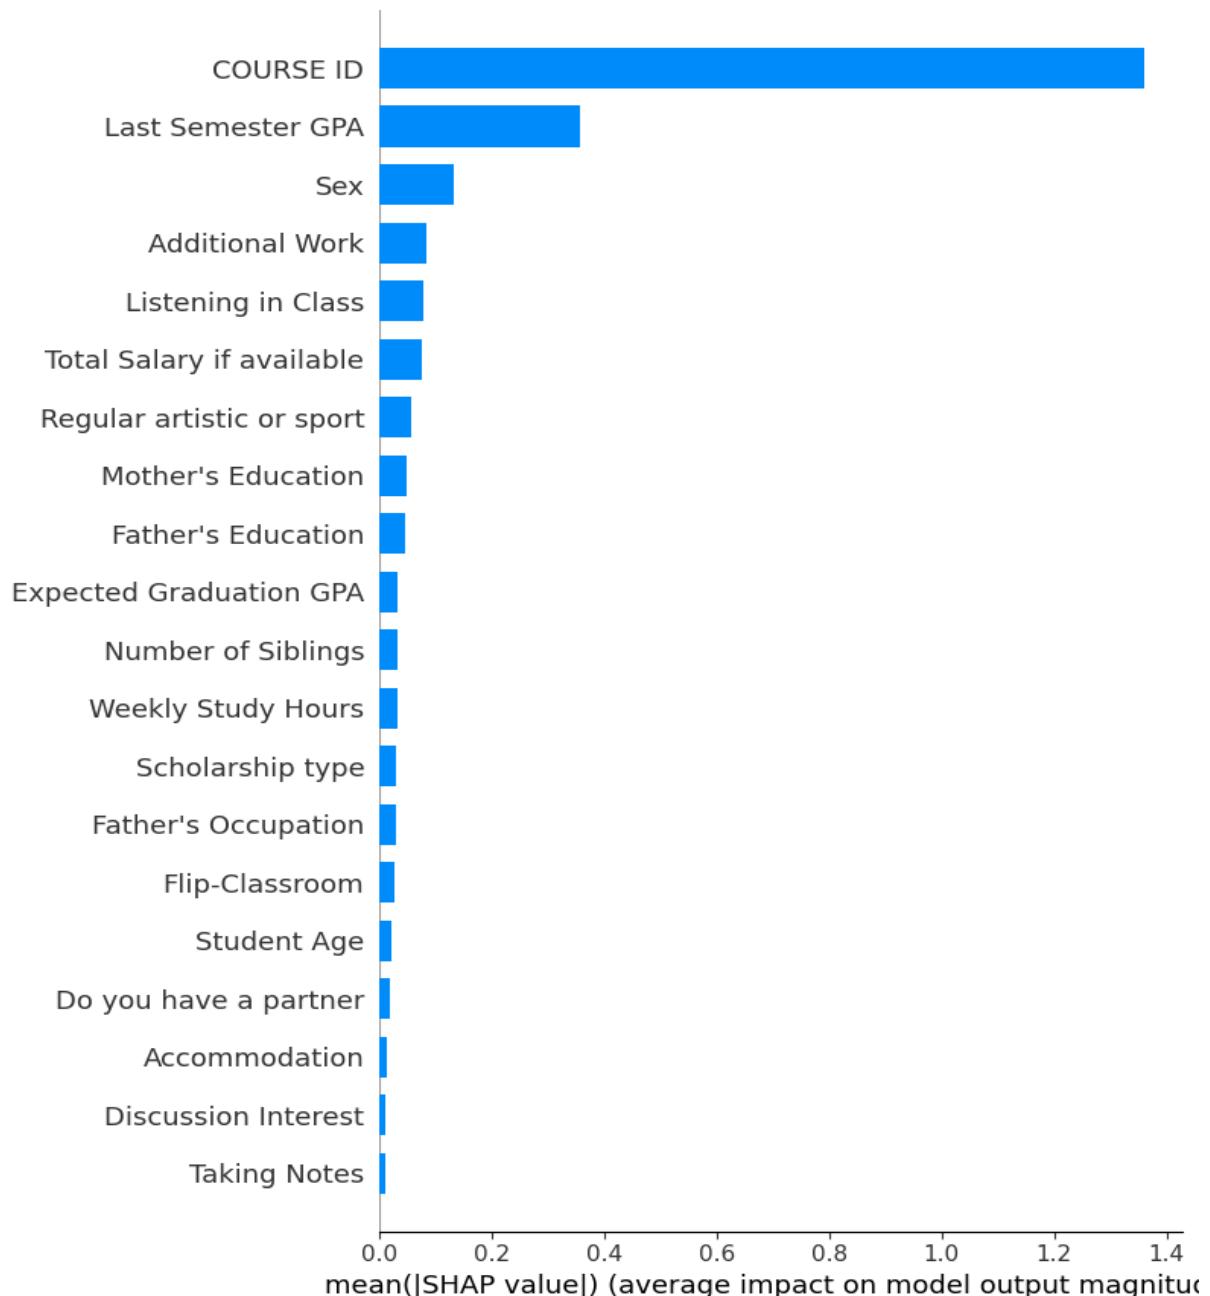
This is an advanced chart that explains direction.

- **Red Dots:** High value (e.g., lots of study hours).
- **Blue Dots:** Low value (e.g., very few study hours).
- **Right Side:** Increases the predicted grade.
- **Left Side:** Decreases the predicted grade.

How to read it:

- Look at **Expected Graduation GPA**. You will see **Red Dots on the Right**.
- **Translation:** "High expectations (Red) lead to higher grades (Right)."
- Look at a negative factor. You might see **Red Dots on the Left**.
- **Translation:** "High values of this feature lower the grade."

8. SHAP Bar Plot (Simple Ranking)



What is this?

This is a simplified version of the previous chart. It ignores the "direction" and just focuses on the magnitude (strength) of the impact.

What it tells us:

It confirms our Feature Importance findings but uses a more sophisticated method (Game Theory) to ensure the ranking is fair. It re-confirms that COURSE ID and Expected GPA are the kings of this dataset.